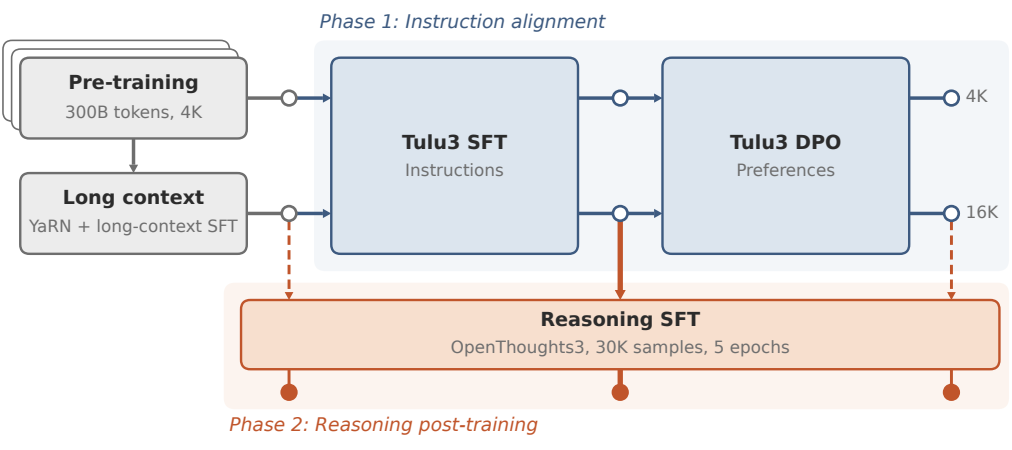
Phase 1: Instruction alignment
Pre-training
300B tokens, 4K
Long context
YaRN + long-context SFT
Tulu3 SFT
Instructions
Tulu3 DPO
Preferences
4K
16K
Reasoning SFT
OpenThoughts3, 30K samples, 5 epochs
Phase 2: Reasoning post-training